

HOME ASSIGNMENT 2, CMPE 260, SPRING 2023. ####
 ####
 #### TOPIC: Sequential decision making with discrete state and action spaces. ####
 ####
 #### LEARNING OBJECTIVES: to get hands-on experience with the theoretical concepts, ####
 #### discussed in the class, such as MDP graph, Value/Policy Iteration, Q-learning, ####
 #### SARSA, DQN (a bonus question). And, to get hands-on experience with efficient ####
 #### implementation with broadcasting of multidimensional tensors. #### ####
 ####
 ####
 #### This assignment is a part of the preparation for the midterm on the same topics. ####
 #### You can work in teams of 2-3 students. You can discuss your solutions with other ####
 #### teams, but sharing your code or parts of it with other teams is plagiarism. ####
 ####
 #### DUE: April 7, 11:59PM. ####

UTILITY FUNCTIONS / PSEUDO-CODE

5 possible actions: Up, Right, Left, Down, Stay
 PSEUDOCODE: Actions = Dict(1=>(-1, 0), 2=>(0, +1), 3=>(0, -1), 4=>(+1, 0), 5=>(0, 0))
 # E.g., Right = (keep constant the first coordinate of state, increase the second coordinate of state by +1)

example for an obstacle (fence with length 3 blocks)
 PSEUDOCODE: Obstacle = [(4, xObst) for xObst in 1:3]

PSEUDOCODE: function validState(s) = # returns true if (state, s, is within the maze boundaries) AND (s is NOT in the obstacles)

dS - size of the maze with dimensions: dS x dS //
 # dA - number of actions
 # Goal - goal state.
 PSEUDOCODE: function BuildMaze(dS, dA, Goal)

dynamics tensor with dimensions: |dS| x |dS| x |dA| x |dS| x |dS| x 1, where the
 # dimensions are S_1, S_2, A, S_1', S_2' . e.g., S_2 is the current second coordinate of the state
 # and S_1' is the first coordinate of the state at the next time step.
 Ps'_sa = zeros(dS, dS, dA, dS, dS)

the reward tensor with the same dimension as the dynamics
 # reward is -1 on every state, and 0 at the Goal state.

```

Rs'sa = -ones(dS, dS, dA, dS, dS)

# iterate over the valid states
for s in filter(validState, (x->x.l).(CartesianIndices((dS, dS))))
    if s ∈ Goal
        Ps'_sa[s..., :, s...] .= 1.0 # all the actions get prob 1 at the goal
        Rs'sa[s..., :, s...] .= 0.0 # all the actions get reward 0
        continue
    end

    for a in Actions # the same action set at each state
        # if "next state is valid" move to it, otherwise stay at place
        s' = validState(s .+ a[2]) ? s .+ a[2] : s
        Ps'_sa[s..., a[1], s'...] = 1.0
    end
end
end
"sanity test:" forall a, s : sum_s' Ps'_sa = 1
return Ps'_sa, Rs'sa
end

```

TASKS

Task 1

Build your maze with dimensions 10x10 and 3 fences, and the goal state (exit) in one of the corners of the maze. Visualize the maze layout on 2D plot.

The tasks 2 - 4 are for the model-based setting, where both $p(s' | s, a)$ and $R(s', s, a)$ are known.

Task 2

Implement the Policy Evaluation (PE) algorithm for a deterministic policy, π .

The dimensions of the policy $\pi[a | s]$ are $|dS| \times |dS| \times |dA| \times 1 \times 1 \times 1$.

There is a single possible action at every state.

E.g., $\pi[\text{UP} | \text{some_state}] = [1, 0, 0, 0, 0]$, see 'Actions' above.

E.g., $\pi[\text{Stay} | \text{another_state}] = [0, 0, 0, 0, 1]$, see 'Actions' above.

Evaluate a random deterministic policy, π . Plot Value of a random policy on 2D plot.

Instructions:

The policy is stationary, which means $\pi[a' | s']$ is `permute_dimensions($\pi[a | s]$, dim1->dim4, dim2->dim5, dim3->dim6)`

Use broadcasting `.*`, e.g.,

$p(s' a' | s, a) = \pi[a' | s'] .* p(s' | s, a)$

$\text{sum}_s p(s' | s, \pi[a | s]) .* V[s']$, where $V[s']$ is the value of the next state with dimensions $1 \times 1 \times 1$

$x \in \mathcal{S} \times \mathcal{S} \times \{1\}$.

The value of the current state has dimensions $d_S \times d_S \times 1 \times 1 \times 1 \times 1$.

"V of the next state" is `permute_dimensions("V of the current state", dims1->dims4, dims2->dims5)`

Task 3

Repeat Task 2 with manually setting the optimal actions in the radius of 2 states from the goal state.

Explain your observations.

Task 4

Implement the Policy Improvement (PI) Algorithm, and find the optimal policy π^* .

Visualize the optimal value function, V_i , on a 2D plot at 3 different iterations, i , of PI.

Explain your observations.

The next tasks are for the model-free setting, where neither $p(s' | s, a)$ nor $R(s', s, a)$ are known.

Task 5

Write a function, `s', r = step(s, a)`, that receives the current state, s , and the current action, a , and

returns the next state, s' , and reward, r .

Generate 10 trajectories from 10 different initial states, using a random policy.

Generate 10 trajectories from 10 different initial states, using the optimal policy from Task 4 above.

Explain your observations.

Task 5

Implement Q-learning algorithm for the tabular case, where Q function is given by a table.

Plot an accumulated reward as a function of the iteration number of Q-learning algorithm for 5 runs of Q-learning from scratch. Plot an average curve of the 4 runs of Q-learning, and the variance (use `fill_between`). Explain your observations

Task 6

Repeat Task 5 with the SARSA Algorithm.

Explain your observations.

Task 7 (bonus 20 points)

Repeat Tasks 5 and 6 with the DQN algorithm, where Q is represented by a neural network with a few

feedforward layers. Compare the results with and without Experience Replay, and with and without a

separate network for the target. Explain your observations.

What to submit:

1. a runnable code, which enables to reproduce your results
 2. a PDF file with all the plots, explanations, and the code.
- Please submit two separate files, rather than a single ZIP, in Canvas.